



Object-Centric Instrumentation with Pharo

Steven Costiou

Square Bracket tutorials

February 20, 2019

Copyright 2017 by Steven Costiou.

The contents of this book are protected under the Creative Commons Attribution-ShareAlike 3.0 Unported license.

You are **free**:

- to **Share**: to copy, distribute and transmit the work,
- to **Remix**: to adapt the work,

Under the following conditions:

Attribution. You must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work).

Share Alike. If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.

For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to this web page:
<http://creativecommons.org/licenses/by-sa/3.0/>

Any of the above conditions can be waived if you get permission from the copyright holder. Nothing in this license impairs or restricts the author's moral rights.



Your fair dealing and other rights are in no way affected by the above. This is a human-readable summary of the Legal Code (the full license):
<http://creativecommons.org/licenses/by-sa/3.0/legalcode>

Contents

Illustrations	iii
1 Introduction	1
1.1 Illustration example	2
1.2 Evaluation criteria	2
1.3 Performance overhead evaluation	3
1.4 Structure of the book	3
2 Summary of the overall evaluations	5
3 Anonymous subclasses	7
3.1 Example	7
3.2 Evaluation	8
3.3 Other documentation	8
4 Talents	9
4.1 Example	9
4.2 Evaluation	10
4.3 Other documentation	10
5 Ghost	11
5.1 Example	11
5.2 Evaluation	11
5.3 Other documentation	12
6 Reflectivity	13
6.1 What are Talents	13
6.2 Example	13
6.3 Evaluation	13
7 Low-level techniques	15
7.1 Example	15
7.2 Evaluation	16
8 Conclusion	17
8.1 Example	17
8.2 Evaluation	18

Bibliography

19

Illustrations

6-1	Installation from Github	13
6-2	Installation from Github	14
7-1	Installation from Github	15
7-2	Installation from Github	15
8-1	Installation from Github	17
8-2	Installation from Github	17



Introduction

This booklet is about object-centric instrumentation in Pharo. An instrumentation is object-centric if it applies to one specific object (or a set of objects), without consideration of its class. It means the instrumentation can be applied on one object, leaving untouched all other instances of its class, or to an heterogeneous set of instances of different classes. This booklet gives an overview of available object-centric instrumentation techniques in Pharo, either present in the standard distribution or available on download. We only focus on object-centric state-access instrumentation, which is a particular case of object-centric instrumentation. We will not go into deep technical usage description, nor into implementation details. Each chapter illustrates one solution with examples, and gives the necessary references if one wants to go deeper in the study of the solution. We study each technique following a three-fold evaluation. First, the studied technique is applied on a simple example of object-centric instrumentation. Second, the technique is evaluated against a set of desirable properties. Finally, performance overhead are evaluated. Only the raw solution is evaluated, without considering the possibility of enhancing the technique by building something on top.

This chapter presents the three-fold evaluation applied to each studied technique, based on the current stable Pharo 7. Each time, a new Pharo image is created, the evaluation code is loaded as well as the studied solution's packages if needed. Then the evaluation is performed. The evaluation code presented in this chapter is available on Github at the following address:

<https://github.com/StevenCostiou/PharoObjectCentricEvaluationExamples>

1.1 Illustration example

Each studied solution is experimented on an example of object-centric behavior instrumentation. We use a class `Person` defined in the following script. This class has a name instance variable and a `name:` method. This method stores the parameter it is given into the instance variable. We would like that each time a value is stored in that instance variable, that value is printed on the Transcript.

```
1 [ Person >> name: aName
2   name := aName
```

The instrumentation can be defined as follows, if `aName` is the reference to the value being stored in the `name` instance variable of the `Person` instance:

```
1 [ aName logCr
```

The evaluation example is defined in the following script. Two instances `p1` and `p2` of class `Person` are created, and object-centric instrumentation must be applied to the `p2` instance. Then each of these instances is given a name through a call to the `name:` method. The result must be that `p2` prints its name in the Transcript, while nothing must happen for `p1`.

```
1 [ |p1 p2|
2   p1 := Person new.
3   p2 := Person new.
4   "Instrumentation must be applied to p2 here"
5   p1 name: 'Worf'.
6   p2 name: 'Dax'.
7   "Only 'Dax' prints in the Transcript"
```

1.2 Evaluation criteria

Each solution is evaluated against the following desirable properties.

Property	Definition
Manipulated entity	The unit of instrumentation (<i>e.g.</i> a class, a Trait, an object...)
Reusability	The entity can be reused to instrument different objects
Flexibility	Instrumentation does not put constraint on the source code or in the coding style
Granularity	The level of at which behavior can be instrumented (<i>e.g.</i> method, AST...)
Integration	Instrumentation does not break system features

1.3 Performance overhead evaluation

To provide a approximation of the performance overhead due to instrumentation, we compare the execution time of a block of code without instrumentation with the execution time of an instrumented block of code. The method `evaluateOverheadFor:` from the following script shows how the average execution time is computed. The parameter is an instance of `Person` that is either not instrumented (*i.e.* to compute the reference execution time used for comparison) or instrumented by one of the studied techniques. The `#name:` message is sent a thousand times to the `Person` instance and each time the execution time is recorded. An average of all the execution times is computed and returned by the method. This average time is used to compare execution time of an instrumented instance against the execution time of a non-instrumented instance.

```

1 evaluateOverheadFor: aPerson
2   |execTimes|
3   execTimes := OrderedCollection new.
4   1 to: 1000 do:[:i|
5     execTimes add: [aPerson name: i] timeToRun].
6   ^execTimes average

```

1.4 Structure of the book

The second chapter will provide an overview of the evaluation results of object-centric instrumentation techniques available in Pharo. A reader may directly read this chapter if he is already familiar with the Pharo techniques presented in the book. Chapters 3 to 7 describe five solutions for object-centric instrumentation, and provide an evaluation of these solutions. Chapter 8 drafts the premises of an object-centric debugger and concludes the book.

CHAPTER 2

Summary of the overall evaluations

If you already know Pharo and (some of) the presented technique, this chapter is a global summary with spoilers.



Anonymous subclasses

Anonymous classes are nameless classes that are inserted between an object and its original class [FJ89, HJJ93]. The object is migrated to that new class, which takes the original object's class as its superclass. Methods from the original class can be redefined and reimplemented in the anonymous class, having the effect to change the behavior of that single object. Original behavior that is not redefined in the anonymous subclass is preserved. It is one of the fastest implementation for object-centric instrumentation [Duc99].

3.1 Example

Anonymous subclasses are derived from the original class of the object (line 3). Methods must be manually (re)written with instrumentation and compiled in the new class (line 4-8). Then the object has to be migrated to its new class (line 10). To rollback the instrumentation, the object must be manually migrated back to its original class (line 12). The migration is not *safe* if more than one process is using the instrumented object.

```
1 |person anonClass|
2   person := Person new.
3   anonClass := anObject class newAnonymousSubclass.
4   anonClass
5     compile:
6       'name: aName
7         self tag: aName.
8         name := aName'.
9   "migrates the object to its new class"
10  anonClass adoptInstance: person.
11  "migrates back the object to its original class"
12  anonClass superclass adoptInstance: person.
```

3.2 Evaluation

Manipulated entity: Classes. Behavioral variations are expressed in standard methods, compiled in anonymous classes.

Reusability: Partial. As an anonymous subclass is derived from the original class of an object, only instances of that same class can be migrated to the anonymous subclass. To apply the same instrumentation to an instance of another class, a new anonymous subclass must be created and the instrumented behavior must be recompiled in that subclass.

Flexibility: None. Instrumented methods must always be copied down to anonymous subclasses, and instrumentation must be inserted in the duplicated code. Without any tool built on top, that instrumentation is fully manual. Anonymous subclasses cannot be composed.

Granularity: Method. Instrumentation is implemented by recompiling modified copies of methods in anonymous subclasses. Sub-method level is achieved through manual rewriting of the method.

Integration: Partial. The object is migrated to an anonymous subclass, which does not break system tools. However, it is explicit that the object is now instance of an anonymous subclass. It may also break libraries and tools that use classes and class names as a discriminator.

Self problem: Solved. By design, `self` always references the original object.

Super problem: Not solved. There are no means to express how to resolve the lookup when a message is sent to `super` from a method copied down in an anonymous subclass.

3.3 Other documentation

The Pharo Mooc provides materials on object-centric instrumentation based on object class migration and its flavours:

- <http://rmod-pharo-mooc.lille.inria.fr/MOOC/Slides/Week7/C019-W7S04-OtherReflective.pdf>
- http://rmod-pharo-mooc.lille.inria.fr/MOOC/WebPortal/co/content_78.html



Talents

Talents are originally behavioral units, that can be attached to an object to add, remove or alter behavior [RGN⁺14]. Only the object to which a talent is attached is affected by behavioral variations. The latest talent implementation relies on trait definition and anonymous subclasses. Talents can be considered as object-centric, stateful-traits.

4.1 Example

Talents are based on traits. Objects can answer to the `#addTalent:` messages (line 9), which takes a `Trait` as parameter. All behavior defined in the trait is flattened in the object. In the following illustration, we instantiate an anonymous trait (line 3), and we compile a method in this trait (line 4-8). That method is an instrumented version of the original `name` method of the class `Person`. This new method replaces the original one, until the talent is removed from the object (line 10). Talents now relies on anonymous subclasses, to which behavior is flattened before objects are migrated to the anonymous class.

```
1 |person talent|
2 |  person := Person new.
3 |  talent := Trait new.
4 |  talent
5 |    compile:
6 |      'name: aName
7 |        self tag: aName.
8 |        name := aName'.
9 |  person addTalent: talent. "adds the talent to the object"
10|  person removeTalent: talent. "removes the talent from the object"
```

4.2 Evaluation

Manipulated entity: Trait. Behavioral variations are expressed using traits. It can be Traits defined in the image or anonymous trait instances in which specific behavior is manually compiled by the developer.

Reusability: Yes. A trait can be added as a Talent to any number of objects.

Flexibility: Partial. Using anonymous traits forces the user to manually compile code in the method. This is however necessary to achieve a sub-method granularity. Conflicts must be resolved manually when Traits are composed.

Granularity: Method. Traits add, remove or alter (through aliasing) the behavior of a method. It can be done at a sub-method level (*e.g.* inserting a statement in the body of a method), but that requires manual rewriting of the method in the Trait.

Integration: Partial. The object is migrated to an anonymous subclass, which does not break system tools. However, it may break libraries that uses classes and class names as a discriminator.

Self problem: Solved. By design, `self` always references the original object.

Super problem: Solved. By flattening all methods that should be found in the super class into the anonymous subclass, and by replacing message sends to super by message sends to `self`.

4.3 Other documentation

The new implementation of Talents is available and documented on Github:

- <https://github.com/tesonep/pharo-talents>



Ghost

Ghost is a general and uniform proxy implementation[PBF⁺15]. A proxy replaces an object to control access to that object [ABW98]. Object-centric instrumentation by means of proxies consist of swapping an object (and all its references) by a proxy object (and references to that proxy object). A proxy object is instance of a proxy class, in which access control is defined. Access control is generally implemented through a single interface, which is called each time a message is intercepted by the proxy. Control behavior then decides what to do with the received message.

5.1 Example

In this example, we use the original implementation of Ghost [PBF⁺15].

5.2 Evaluation

Manipulated entity: Classes. Proxies are defined and/or configured in classes which inherits from Ghost internal classes. Typically, developers subclass the base message handler from Ghost to create a proxy model that implements the wanted instrumentation. An API is provided to apply a proxy to objects.

Reusability: Complete. The same proxy model can be reused to instrument any kind of object with the same instrumentation. Although Ghost proxies are not meant to be composed, it should be possible to *proxify* a proxy, if a proper model is implemented to instrument a proxy object.

Flexibility: Partial. Because the user has to define a proxy model which specifies which messages are handled and/or which are not. However that

allows developers to fully and transparently integrate the instrumented object into the environment.

Granularity: Method. A proxy intercepts messages sends to the object it *proxifies*. It can execute instrumentation behavior before, after or instead the intercepted message. Sub-method instrumentation cannot be achieved by means of proxies.

Integration: Full. because meat-messages can bde efined

Self problem: Implementation dependent.

Super problem: Implementation dependent. ?

5.3 Other documentation

Implementations and documentation based on the original Ghost paper [PBF⁺15]:

- http://esug.org/data/ESUG2011/IWST/PRESENTATIONS/23.Mariano_Peck-Ghost-ESUG2011.pdf
- <https://rmod.inria.fr/archives/papers/Mart14z-Ghost-Final.pdf>
- <https://gitlab.inria.fr/RMOD/Ghost>
- <https://github.com/guilleg/avatar>

Another implementation of Ghost:

- <https://github.com/pharo-ide/Ghost>
- <http://dionisiydk.blogspot.com/2016/04/halt-next-object-call.html>



Reflectivity

Talents [RGN⁺14]is this.

6.1 What are Talents

6.2 Example

Installing Talents

aa

Example

bb

6.3 Evaluation

cc

Listing 6-1 Installation from Github

```
1 [
2   Metacello new
3     baseline: 'Talents';
4     repository: 'github://tesonep/pharo-talents/src';
5     load.
```

Listing 6-2 Installation from Github

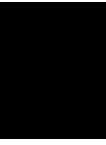
```

1
2 talent := Trait named: 'MyTalent'.
3 talent compile: 'add: anObject
4 anObject logCr.
5 super add: anObject'.
6 col := OrderedCollection new.
7 col addTalent: talent.
8 col add: 'This is an added object.'
```

■ **Note** this is a note annotation.

■ **To do** this is a todo annotation

Country	Capital
France	Paris
Belgium	Brussels
Country	Capital
France	Paris
Belgium	Brussels



Low-level techniques

7.1 Example

Installing Talents

aa

Example

bb

Listing 7-1 Installation from Github

```
1 [
2   Metacello new
3     baseline: 'Talents';
4     repository: 'github://tesonep/pharo-talents/src';
5     load.
```

Listing 7-2 Installation from Github

```
1 [
2   talent := Trait named: 'MyTalent'.
3   talent compile: 'add: anObject
4   anObject logCr.
5   super add: anObject'.
6   col := OrderedCollection new.
7   col addTalent: talent.
8   col add: 'This is an added object.'
```

7.2 Evaluation

cc

■ **Note** this is a note annotation.

■ **To do** this is a todo annotation

Country	Capital
France	Paris
Belgium	Brussels
Country	Capital
France	Paris
Belgium	Brussels



Conclusion

8.1 Example

Installing Talents

aa

Example

bb

Listing 8-1 Installation from Github

```
1 [
2   Metacello new
3     baseline: 'Talents';
4     repository: 'github://tesonep/pharo-talents/src';
5     load.
```

Listing 8-2 Installation from Github

```
1 [
2   talent := Trait named: 'MyTalent'.
3   talent compile: 'add: anObject
4   anObject logCr.
5   super add: anObject'.
6   col := OrderedCollection new.
7   col addTalent: talent.
8   col add: 'This is an added object.'
```

8.2 Evaluation

cc

- **Note** this is a note annotation.
- **To do** this is a todo annotation

Country	Capital
France	Paris
Belgium	Brussels
Country	Capital
France	Paris
Belgium	Brussels

Bibliography

- [ABW98] Sherman R Alpert, Kyle Brown, and Bobby Woolf. *The design patterns Smalltalk companion*. Addison-Wesley Longman Publishing Co., Inc., 1998.
- [Duc99] Stéphane Ducasse. Evaluating message passing control techniques in smalltalk. *Journal of Object Oriented Programming*, 12:39–50, 1999.
- [FJ89] Brian Foote and Ralph E Johnson. Reflective facilities in smalltalk-80. In *ACM Sigplan Notices*, volume 24, pages 327–335. ACM, 1989.
- [HJJ93] Bob Hinkle, Vicki Jones, and Ralph E Johnson. Debugging objects. In *The Smalltalk Report*. Citeseer, 1993.
- [PBF⁺15] Mariano Martinez Peck, Noury Bouraqadi, Luc Fabresse, Marcus Denker, and Camille Teruel. Ghost: A uniform and general-purpose proxy implementation. *Journal of Object Technology*, 98:339–359, 2015.
- [RGN⁺14] Jorge Ressia, Tudor Gîrba, Oscar Nierstrasz, Fabrizio Perin, and Lukas Renggli. Talents: an environment for dynamically composing units of reuse. *Software: Practice and Experience*, 44(4):413–432, 2014.

